



AI for Competitive Play in Terra Mystical

Luis Antonio Perez

Computer Science Department, Stanford University

CS 221: Artificial Intelligence

Stanford
University

Motivations and Goal

Problem and Goal

- Create an AI agent capable of playing TM at super-human level.

What is Terra Mystical (TM)

- 2-5 player strategic board game with collaborative components.
- Each player tries to maximize the number of victory points (VP) earned during six fix rounds of gameplay,

Why TM

- Recent breakthroughs in strategy games (Go and Chess/Shogi)
- Complex gameplay mechanics, multiple “actions” per game
- No current AI has beaten top human-players



Monte Carlo Tree Search

SELECTION


Traverse Existing Search Tree with Deterministic Policy Guided by p

EXPANSION + EVALUATION

On non-terminal leaf-node, expand using NN to compute $p(s)$ and $v(s)$

BACKPROPAGATION

Use estimated $v(s)$ for leaf-state and propagate this result to all $Q(s,a)$ of path




Self-Play Reinforcement Learning

- Value Estimation** - for newly encountered states, estimate value and use this as proxy for full MCTS rollout
- Neural Network Guided MCTS search** - prior probabilities set by NN predictions
- Loss Function**
 - Collect set of \mathbf{p}_i , \mathbf{z} and use to update NN predictions
 - Can occur while performing self-play - parameters updated immediately

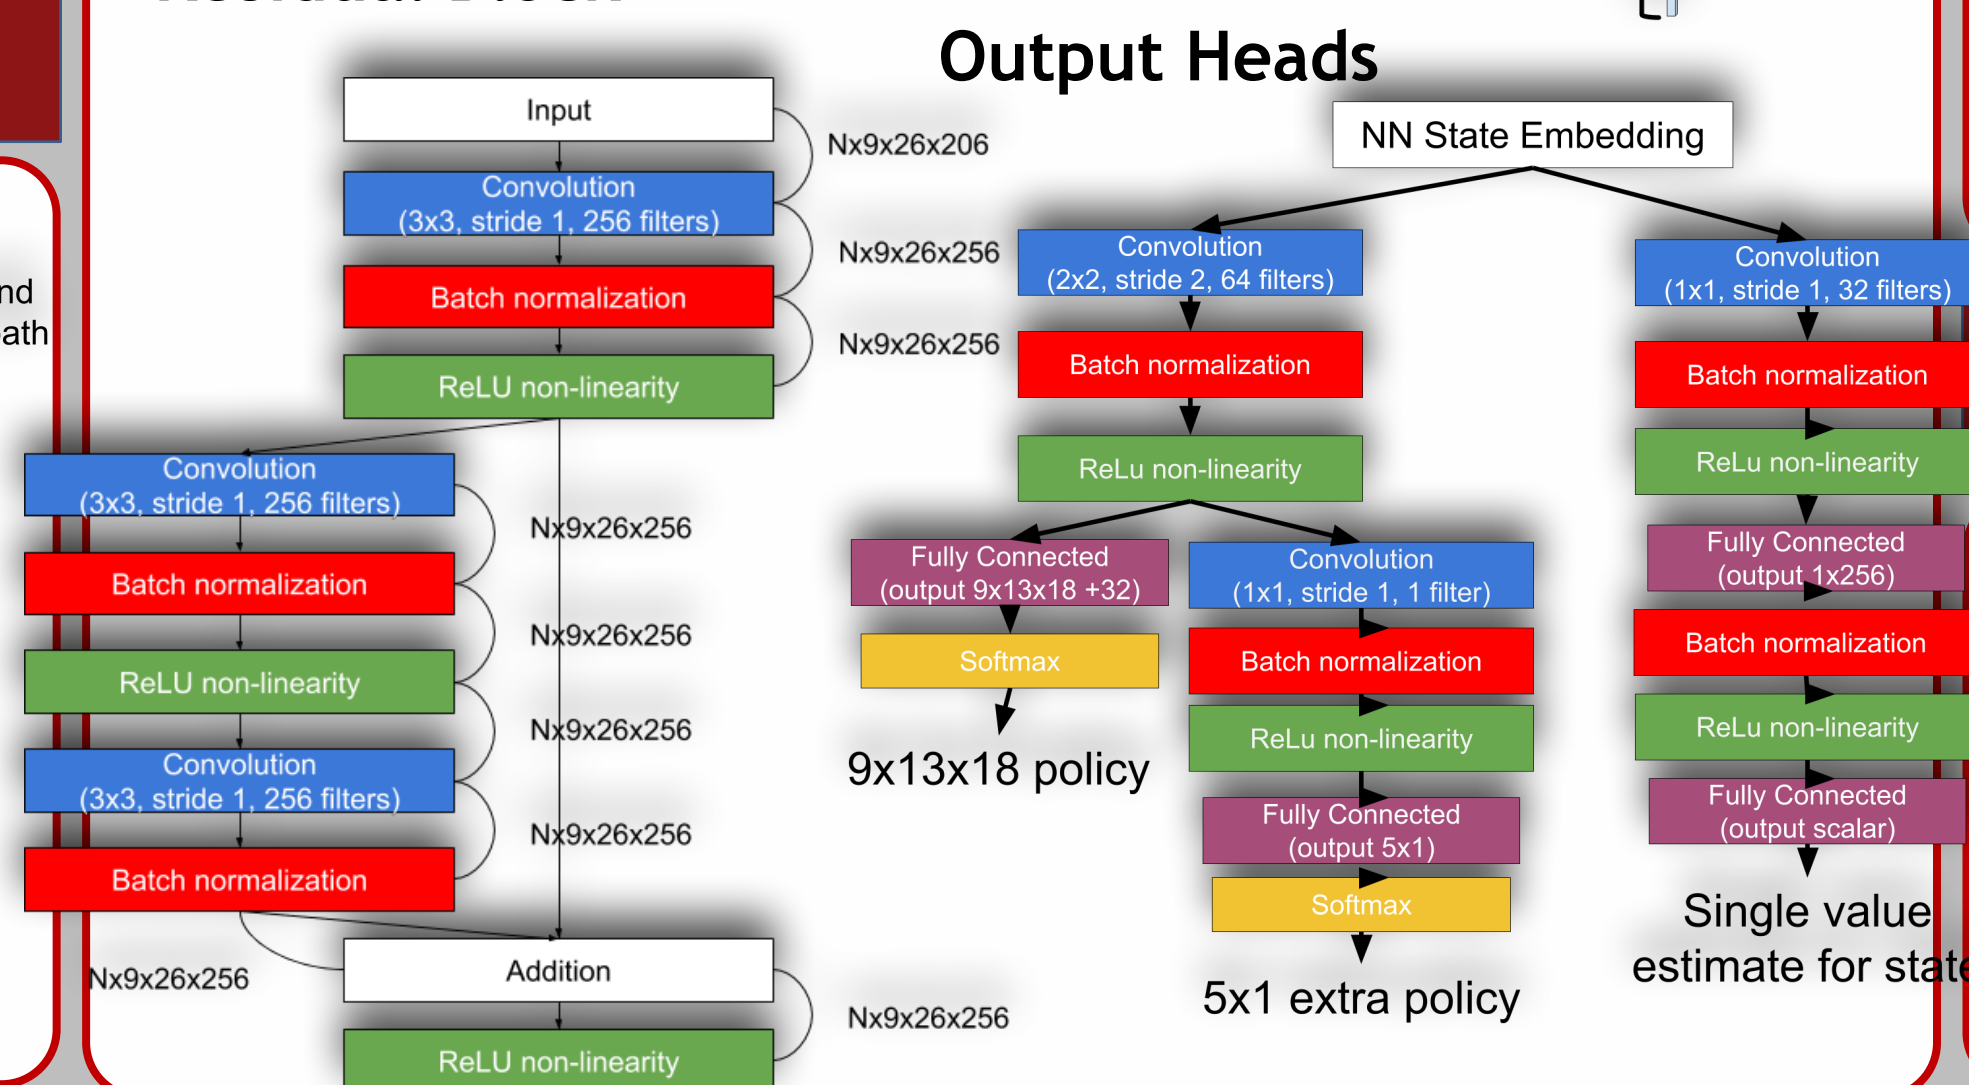
$$(\mathbf{p}, \mathbf{v}) = f_{\theta}(\mathbf{s}) \text{ and } l = (\mathbf{z} - \mathbf{v})^2 - \pi^T \log \mathbf{p} + c \|\theta\|^2$$

Neural Network Architecture

Multi-headed Architecture



Residual Block



Results

Player Faction/AI Algorithm	Human Player	State of the Art AI	Alpha TM
Halflings	133.32	92.21	32.11
Engineers	127.72	77.12	34.12

Future Work

Introduce Supervised Learning to Stabilize Training

- Difficulty training using only self-play
- There are XX,000 games available online among human and existing AI player - use this information to jump-start self-play, similar to AlphaGo.

Expand to Multi-Player

- Current work focuses on 2-player games
- Next steps is to expand self-play to multi-player

Introduce Different NN Architectures

- Introduce better policy and value estimators
- Possible use for RNNs or LSTMs for long-term memory storage and strategy

References and Acknowledgements

- [1] BoardGeek. Terra mistica: Statistics, 2011
 - [2] D. Silver, et. al, Mastering the game of go with deep neural networks and tree search. Nature, 529:484 EP –, Jan 2016. Article.
 - [3] D. Silver, et. al, Mastering chess and shogi by self-play with a general reinforcement learning algorithm. CoRR, abs/1712.01815, 2017.
 - [4] D. Silver, et.al. Mastering the game of go with- out human knowledge. Nature, 550:354 EP –, Oct 2017. Article.
 - [5] tmai. Terra mystic ai implementations, 2018
- We would like to acknowledge Professor Liang and our mentor Steven Diamond